

# Sparbit: a new logarithmic-cost and data locality-aware MPI Allgather algorithm

Wilton Jaciel Loch

*Graduate Program in Applied Computing  
Santa Catarina State University (UDESC)  
wilton.loch@edu.udesc.br*

Guilherme Piégas Koslovski

*Graduate Program in Applied Computing  
Santa Catarina State University (UDESC)  
guilherme.koslovski@udesc.br*

**Abstract**—The collective operations are considered critical for improving the performance of exascale-ready and high-performance computing applications. On this paper we focus on the Message-Passing Interface (MPI) Allgather many to many collective, which is amongst the most called and time-consuming operations. Each MPI algorithm for this call suffers from different operational and performance limitations, that might include only working for restricted cases, requiring linear amounts of communication steps with the growth in number of processes, memory copies and shifts to assure correct data organization, and non-local data exchange patterns, most of which negatively contribute to the total operation time. All these characteristics create an environment where there is no algorithm which is the best for all cases and this consequently implies that careful choices of alternatives must be made to execute the call. Considering such aspects, we propose the Stripe Parallel Binomial Trees (Sparbit) algorithm, which has optimal latency and bandwidth time costs with no usage restrictions. It also maintains a much more local communication pattern that minimizes the delays due to long range exchanges, allowing the extraction of more performance from current systems when compared with asymptotically equivalent alternatives. On its best scenario, Sparbit surpassed the traditional MPI algorithms on 46.43% of test cases with mean (median) improvements of 34.7% (26.16%) and highest reaching 84.16%.

**Keywords**-Message-Passing, MPI, Collective Communication, Collective Algorithms, Allgather.

## I. INTRODUCTION

The Message Passing Interface (MPI) standard defines several constructs to handle communication between concurrent processes. These can be generally divided in two forms of message exchanges: point to point operations, where two processes communicate individually with each other through one- or two-way data movements; and collective operations, where multiple processes contribute as data sources, sinks or both into a broader coordinated exchange [1]. Although both forms have significant usage on High-Performance Computing (HPC) applications, the latter is responsible for the majority of communication time spent on MPI and also accounts for the most number of calls [2]. Therefore, the collective operations are considered critical to improve the performance of parallel and distributed memory systems [3]. On this paper, we focus on the Allgather (`MPI_Allgather`) many to many operation, which amongst other collectives

holds a large share of utilization [4] and time consumption [2] on applications.

The delay magnitude for the completion of a collective is a product of several factors, which include the underlying hardware topology [5], communication protocols, network capacity [3], placement of processes [6], [7] and others. However, one of paramount importance is the performance of the algorithm employed to coordinate the high level inter process communication and block transfers [8], [9]. On a purely theoretical regard, the community has long discovered the minimum time costs to perform this task, both in the latency term - or the required number of steps - and in the bandwidth term - or the actual data transfer time [10]. This, in turn, led to the development of theoretically optimal algorithms which are capable of performing all the exchanges while maintaining minimum costs.

For Allgather, the generally available algorithms are Ring, Neighbor Exchange, Bruck and Recursive Doubling. The Ring algorithm has a linear growth of both latency and bandwidth time with the increase in the number of processes [7]. Neighbor Exchange has the same asymptotic behaviour but with a less steep increase in time, requiring only half of the Ring's steps, with the downside of only working for even numbers of processes [11]. Recursive Doubling reaches the minimum costs by delivering the data with a logarithmic increase in latency, but only works with power of two numbers of processes [12]. Finally, Bruck works for any number of processes and also has a logarithmic cost, requiring only one additional step in non power of two process numbers [10].

One could expect that for being optimal cost algorithms both Recursive Doubling and Bruck would be employed at all cases. Nonetheless, the logarithmic options may not yield the best performances when compared with theoretically less efficient ones as Ring [12] and Neighbor Exchange [11], and the reason for these observations lies on practical limitations that arise when the algorithms are implemented. In practice, Bruck and Recursive Doubling are only employed for reduced data sizes, because although both have the smallest costs, their communication patterns are too diffuse, which severely hurts their performance with large data sizes. As the size of the block to be transferred grows, Ring and Neighbor Exchange are employed because

despite having a worse cost, their communication patterns are local and hence, less expensive to move data [12]. Figure 1 presents a heat map displaying the algorithm with smallest time for each combination of data size and number of processes (experiments are further detailed on Section IV) and portrays the aforementioned characteristic of performance degradation for logarithmic algorithms on large data sizes, where linear options achieve the best times. It is also clear how on power of two numbers of processes Recursive Doubling generally achieves better times than Bruck, while a similar pattern can be seen for even numbers of processes with Neighbor Exchange in relation to Ring. All these observations highlight the fact that there is no *silver bullet* algorithm which is the best for all scenarios and test cases.

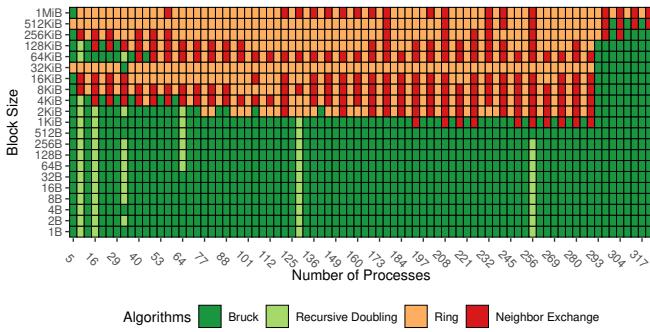


Figure 1: Algorithms with the best experimental time for each case.

On the current collective algorithm scenario, although certain options provide the minimum theoretical costs, their performance is damaged by non-local communication patterns. This hindering phenomenon is so severe that it turns linear algorithms into more efficient alternatives, rendering the advantage of logarithmic number of steps virtually useless in many cases. We therefore propose the Stripe Parallel Binomial Trees (Sparbit) algorithm, which has logarithmic costs and still maintains local data exchange patterns. Sparbit was better than the currently employed algorithms on 46.43% of test cases (624/1344) with mean (median) improvements of 34.7% (26.16%) and highest reaching 84.16% on its best scenario.

The additional performance provided by our novel algorithm, which improved upon the state of the art in several cases, will only be leveraged if an efficient algorithm selection policy is employed. This however has by itself a consistent dedicated research effort, like the works of [13] and [14] as well as the Open Tool for Parameter Optimization (OTPO) focused on Open MPI. Therefore, we retain our focus towards presenting and validating the proposal, while usage guidelines are the aim of other independent works. This paper is organized as follows. Section II provides a deeper look into the currently employed Allgather algorithms. Sparbit and its functioning are introduced on Section

III. The experimental methodology is provided on Section IV alongside benchmarks, configurations and other important information. On Section V the experimental results are analysed, Section VI presents the related work and Section VII concludes the paper.

## II. ALGORITHMS AND PROBLEM DEFINITION

The MPI\_Allgather call is a many to many collective operation where each process sends a fixed size block of data to all other participants and receives their individual blocks as well. The block sent by a process has one or more items of a chosen datatype and both the size and type signature of the block must be equal to all destinations. After the Allgather execution, the block sent by the  $j$ th process is received by all others and placed in the  $j$ th position of their receive buffer [1], whose final configuration will be equal to all processes. Next, we provide a brief description of the algorithms gathered from the literature and from source code analysis of both MPICH 3.3.2 and Open MPI 4.0.3. Our focus was directed towards these implementations since they are open source and the base for other closed source ones. Then, driving from the presented information and discussion, we formulate the problem.

### A. Allgather algorithms

Initially, a cost model is necessary in order to accurately compare the algorithms that implement the Allgather operation. To this end we present each algorithm's complexity as cost in time following the Hockney model [15]. With this approach the cost for sending a message is given by a latency or start-up term  $\alpha$  and a bandwidth or transference cost per byte term  $\beta$ . Since all the presented algorithms essentially implement communication patterns built of individual point to point messages, the total cost of each algorithm following the Hockney model is given by the time sum of all the building block messages.

**Ring:** The Ring algorithm performs the block exchanges by shifting data around on a ring communication topology. On every step, a process with rank  $r$  will send a block to the process with rank  $r+1$  and receive another from the process with rank  $r-1$  (wrapping around if a destination or source is out of bounds). Each process's own block is sent on the first step, while on all others the block received on the previous step is forwarded [7]. Hereafter, the number of processes involved in the algorithm is represented by  $p$ , while  $m$  represents the total amount of data that a process must have at the end of the operation. In this sense, each process has to receive a block with size  $\frac{m}{p}$  from all the other  $p-1$  participants as well as send its own. Since all sends happen in parallel and any block takes  $p-1$  steps to reach all processes, the cost is given by  $C_{ring} = (p-1)(\alpha + \frac{m}{p}\beta)$ . The ring is employed both on MPICH and Open MPI for various cases.

**Neighbor Exchange:** The Neighbor Exchange is another linear complexity algorithm, however its latency growth is smaller than the Ring by half. Its main driving idea is to make all transferences pairwise to better leverage the TCP protocol, such that each process on a step  $s$ , with  $0 \leq s < \frac{p}{2}$ , will only communicate with one peer, both sending and receiving data. Following this, an even process  $r$  will send data to its neighbor  $r + (-1)^s$  while an odd process  $r'$  will send data to its neighbor  $r' - (-1)^s$ , both with wrap-around. As two blocks are sent on all steps but the first, this algorithm's cost given by  $C_{ne} = \frac{p\alpha}{2} + (p-1)\frac{m}{p}\beta$  [11]. The negative side is that it only works with an even number of processes. The Neighbor Exchange is not employed on MPICH, but is utilized for several cases on Open MPI.

**Recursive Doubling:** The third available algorithm is the Recursive Doubling, which performs the complete exchange in a logarithmic amount of steps. The main difference of this algorithm in relation to the previous ones is that it doubles both the distance of the communication and the data being sent. On every step  $s$ , with  $0 \leq s < \log_2 p$  a process with rank  $r$  will exchange data in a pairwise fashion with a process that has rank  $r \oplus 2^s$ , where  $\oplus$  represents the binary exclusive *or*. As all process send all the data received so far, the number of blocks double at every step and the cost of the Recursive Doubling is given by  $C_{rd} = (\log_2 p)\alpha + (p-1)\frac{m}{p}\beta$  [16]. The operation of this algorithm is limited only to numbers of processes that are powers of two and thus this is the only case where it is employed both on MPICH and Open MPI.

**Bruck:** The Bruck is very similar to the Recursive Doubling on its core, doubling distances and data. However, it does not employ pairwise exchanges and on every step  $s$ , with  $0 \leq s < \lfloor \log_2 p \rfloor$ , a process  $r$  will send all the data it has to the process with rank  $r - 2^s$  and receive new data from the process with rank  $r + 2^s$ . If the number of processes is not a power of two, an additional step is needed where only the first  $p - 2^{\lfloor \log_2 p \rfloor}$  blocks in the receive buffer are sent [10]. The cost of the Bruck is given by  $C_{bruck} = \lceil \log_2 p \rceil \alpha + (p-1)\frac{m}{p}\beta$ , it has no usage restrictions and is employed both on MPICH and Open MPI for various scenarios.

### B. Problem formulation

So far while presenting the algorithms we have only discussed usage limitations that are inherent to their architecture. However, in practice the algorithms also have usage limitations that stem from performance issues, which tend to be side effects of their relation with the supporting technologies. The formal definitions of the algorithms assume equally balanced communication costs to all peers, but computing clusters and supercomputers often employ hierarchical network topologies [6]. On these networks the cost for performing communication between two nodes is highly dependent on the physical location of each peer [7], and the further away they are, the longer are the physical

paths between them and therefore the higher the latency. From a bandwidth perspective, the further away two nodes are the higher is the chance that their communication will cross the core of the network, whose bandwidth is more expensive and supports less saturation than the edge [17], possibly leading to slowdowns or contentions.

Along these lines, although Recursive Doubling and Bruck have a smaller number of steps, they naturally communicate over longer paths on the network, implying in higher time costs specially for larger messages. Additionally, Bruck also needs memory shifts to assure contiguous sends and data organization, which further deteriorate its performance with the increase in data size. On the other hand, Ring and Neighbor Exchange have worse linear latency costs, but have local communication patterns always with the same neighbors, which implies in reduced additional transfer costs with the growth of message size [12].

This scenario outlines a conflict between the theoretical quality of the algorithms and their actual efficiency: so far, the performance of logarithmic algorithms is degraded by their non-local communication to such an extent that they achieve worse times than linear options in many cases. This phenomenon hinders their potential for reduced time costs and limits the overall application performance.

## III. THE SPARBIT ALGORITHM

This section outlines the design of the Stripe Parallel Binomial Trees (Sparbit) algorithm. It is inherent to the logarithmic communication algorithms that the amount of data exchanged at any given step will be larger than the amount of the previous one, usually by a factor of two, implying that the cost of data exchange grows at every step as a product of data size. This is true for Bruck, Recursive Doubling and syllogistically for Sparbit too. Additionally, if we suppose the usage of hierarchical network topologies, then the cost of sending and receiving data increases with the growth of distance between any two communicating peers. We can hence derive a cost estimate for every step as being a product of these two dimensions: the larger the data size and the further the distances travelled by the data, the higher will be the cost in time for the communication. The former dimension can not be changed in order to maintain the logarithmic complexity, however there is no restriction on the latter. Stemming from this and through a higher behaviour perspective, the goal of Sparbit is to provide the same data exchange semantics as the Bruck algorithm, however instead of doubling the distances starting from 1 until  $2^{\lceil \log_2 p \rceil - 1}$ , it takes the opposite direction and halves them starting from  $2^{\lceil \log_2 p \rceil - 1}$  until 1. This creates a more balanced distribution of communication costs along the execution, since as the data sizes inevitably grow, the distance they must traverse continuously shrinks.

### A. Binomial tree

As suggested on its name, Sparbit uses the binomial tree to deliver data, which is a one to many logarithmic data

**delivery algorithm.** Initially on its execution (Figure 2), the block of data resides in the root process, here assumed to have rank 0. On the first step, the root will send the block to the process with rank  $p/2$  (P4) and will delegate to it the responsibility of delivering the data to the processes with rank  $r > p/2$ . Therefore, this first step has divided the group of processes in two subtrees, with each root 0 and  $p/2$  responsible for sending the block to half of the processes. The algorithm is then repeated recursively and in parallel for each of the halves, continuously subdividing the processes into smaller subtrees until the data has reached all of them.

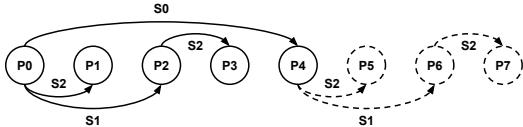


Figure 2: Binomial tree. Dashed circles and arrows represent missing processes and ignored sends, respectively.

For a non power of two number of processes, as exemplified with 5 on Figure 2, the algorithm works by exactly mimicking the way a binomial tree with size  $2^{\lceil \log_2 p \rceil}$  would work, however ignoring any sends that are destined to processes with ranks equal or greater than  $p$ . Since all sends happen in parallel, the cost of the binomial tree algorithm is given by  $C_{\text{binomial\_tree}} = \lceil \log_2 p \rceil(\alpha + m\beta)$ .

### B. Algorithm design

Regarding the semantics of data delivery on an MPI All-gather call with  $p$  participants, each process must somehow send its block to all the other  $p - 1$  peers. If we take only the process with rank 0 and suppose that it is the root of a binomial tree composed of all the participating processes, then it would be trivial to deliver its block to all the destinations in a logarithmic number of steps. Moreover, if we imagine that the process rank space is circularly connected (*i.e.*, there is wrap-around on the destination calculation), then we can expand our previous supposition to any process  $r$ , since the tree organization would be exactly the same only shifted by  $r$  units. Finally from this concept, if all processes act as roots of the binomial tree destined to distribute their own block and simultaneously as leafs or intermediate nodes of all the other  $p - 1$  trees, the whole Allgather semantic can be achieved. This idea is the core of Sparbit and virtually all that it does is start and manage the participation of each process in the execution of all these parallel binomial trees. The resulting communication pattern of Sparbit can be also modelled as a Binomial Graph [18], however utilizing a unique data distribution behaviour. Bruck employs a similar general concept but with inverse distance circulant graphs, which manifest themselves in the form of parallel spanning trees that cause its locality issues.

Although the idea of Sparbit is simple, there is a considerable disconnect between its observed behaviour and the actual algorithmic logic that makes it work. On the execution

of the binomial tree distribution, on each step with distance  $d$  a process that already has the block will send it to a peer  $d$  away on the process rank space. If the number of processes is a power of two, then the execution is trivial and on every step  $s$ , with  $0 \leq s < \log_2 p$ , a process both sends and receives  $2^s$  blocks to and from destinations  $2^{\log_2 p - (s+1)}$  away on each direction. However, if the amount of processes is not a power of two, then some block sends will be ignored to avoid unexisting destinations. When employing wrap-around there are no unexisting destinations, but facing the same execution scenario and not ignoring these sends would imply in double writing the same block, which is not desirable. Hence, if we take again the example shown in Figure 2 with 5 processes and imagine that all the blocks initially reside on process 0, then on the first step process 4 would receive them. Now, since process 4 is a leaf and thus has no destinations, it would simply ignore the regular behaviour of forwarding the received blocks. When looking to the real Sparbit scenario the same logic applies, however the blocks are scattered along all the processes, with each distributing its portion through its own individual tree. Since each tree is shifted towards a unique root and all processes participate in all trees, then each process must act as the relative fourth destination -  $(\text{root} + 4)\%p$  - to one of these trees. This implies that the role of ignoring sends, done in the initial supposition only by process 4, now must be done by all processes, each ignoring the blocks received through the trees in which it is a leaf, but still forwarding the ones from trees in which it is an intermediate node.

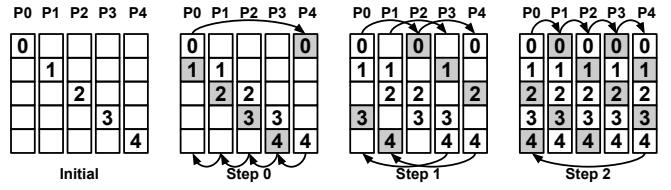


Figure 3: Sparbit buffer configuration with 5 processes on different steps. Grey blocks are received on the current step.

The collective ignoring and forwarding behaviour can be clearly observed on Figure 3, that shows the buffer configuration for each process on different steps of Sparbit's execution. On step 0, each process  $r$  both sends its original block to process  $r + 4$  and receives the original block of process  $r - 4$ , with wrap-around. Since the received block comes from the left, it is placed on the corresponding upwards distance on the receive buffer, also with wrap-around as can be seen on processes 0 to 3. On the beginning of step 1 all processes have two blocks, however since the one received on the previous step comes from trees where the processes are leafs, it must not be forwarded and again only the original block is sent on this step with distance 2. On the last step, each process has 3 blocks and the same one must still be ignored. Nonetheless, the other two belong respectively to the tree in which the process is a root and to

the one where it is an intermediate node, thus both must be forwarded with distance one.

### Algorithm 1: Sparbit.

**input:** The amount of processes  $p$ ; the rank of the current process  $rank$

```

1  data ← 1;
2  ignore ← 0;
3  d ←  $2^{\lceil \log_2 p \rceil} - 1$ ;
4  last_ignore ← count_trailing_zeros(p);
5  ignore_steps ← ( $\neg(p >> \text{last\_ignore}) \mid 1$ ) <<
   last_ignore;
6  for  $0 \leq i < \lceil \log_2 p \rceil$  do
7    if  $d \& \text{ignore\_vector}$  then
8      | ignore ← 1;
9    end
10   for  $0 \leq j < \text{data} - \text{ignore}$  do
11     | MPI_Isend(rank +  $d \% p$ ,
12       |  $(rank - (2j)d + p) \% p$ , 1);
13     | MPI_Recv((rank -  $d + p) \% p$ ,
14       |  $(rank - (2j + 1)d + p) \% p$ , 1);
15   end
16   MPI_Waitall()
17   d ←  $d >> 1$ ;
18   data ← (data << 1) - ignore;
19   ignore ← 0;
20 end

```

Another way to see a binomial tree is as a recursive structure composed of smaller power of two binomial trees, whose roots must receive the data to further distribute it amongst their branches. The amount and size of subtrees is respectively defined by the number and values of the powers of two which compose  $p$ . In the example of Figure 2 with 5 processes there are two such subtrees, one with 4 processes of ranks 0 to 3, and another with 1 process of rank 4. Therefore if we employ this rationale on the example, the first step takes the data from the root of the first tree to the root of the second one, performing what we call a tree expansion. If the number of processes was larger there would be more subtrees, and therefore more tree expansions during the execution. For example, if there were 21 processes, there would be three subtrees of sizes 16, 4 and 1. Also since the distance halves, the trees are organized in descending order, with larger subtrees comprising the smaller relative ranks and the global root always being the root of the first one. Therefore, the block flows from the larger to the smaller subtrees and a tree expansion will only happen on a step with distance  $d$  if there is a subtree of size  $s_1 = d$  that has the block and another subtree with size  $s_2 < s_1$  that does not. Practically, a tree expansion only indicates that on that step every process should send all the blocks it has.

With all this in mind, Algorithm 1 presents Sparbit's pseudocode. Function signatures for `Isend` and `Irecv`

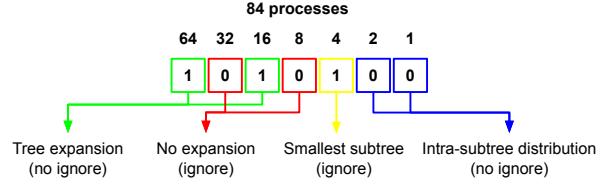


Figure 4: Binary representation with behaviour on each bit.

were simplified for brevity, taking destination/source, buffer displacement in blocks and number of blocks as parameters. The expected amount of data and whether the sending of a block should be ignored on the current step are kept in the `data` and `ignore` variables, respectively. The `ignore_steps` variable stores the steps in which a block should be ignored, being built from the binary representation of  $p$  as shown on Figure 4 for an arbitrary number. For  $t$  subtrees, since the data is already initially located on the first one, only  $t-1$  expansions are needed. These expansions happen on the  $t-1$  larger subtrees and thus no expansion happens on the smallest one. Below the smallest subtree all blocks should be sent on every step as now all distributions are local. Such patterns are highlighted on Figure 4 and if we invert all bits to the left of the first one set, then steps in which a block must be ignored will have their respective bits set, while unset bits will indicate steps with the opposite behaviour. All these binary operations are performed on lines 4 and 5 of Algorithm 1. Through  $\log_2 p$  steps (lines 6-18) the proper blocks get sends and receives issued (lines 10-13) and waited for (line 14), the distance halves as the data size doubles (lines 15 and 16). If in the current step a block must be ignored, the proper variable is set (lines 7-9) and the number of blocks sent is reduced (line 10). Since all sends happen in parallel and Sparbit takes a logarithmic number of steps, its cost is given by  $C_{sparbit} = \lceil \log_2 p \rceil \alpha + (p-1) \frac{m}{p} \beta$ .

### C. Process mappings, topology and Sparbit

The mapping of MPI processes to the topology is done by assigning ranks to their instantiations on the machines. Naturally, since the communications are established by means of process' ranks, the manner in which they are distributed plays a significant role on the respective costs. The two main predefined ways in which MPI implementations map ranks to nodes are *sequential* and *cyclic*. The first performs a best-fit approach, completely filling the slots of a machine with sequential ranks before going to the next. The second uses a round-robin approach, assigning only one rank to a machine and then moving to the next in a circular fashion until they are all filled. By default, Open MPI employs sequential mapping while MPICH employs cyclic [7].

Sparbit is expected to achieve greater locality under sequential mapping, as the heavier communication stages are held among closer ranks, which in this case are physically mapped near to each other. Meanwhile, Bruck is expected to perform better under cyclic mapping due to the same reason.

#### IV. EXPERIMENTAL METHODOLOGY

Two sets of experiments were executed on the Yahoo and Cervino clusters from the University of Neuchâtel, where 16 and 5 dedicated homogeneous machines were employed, respectively. On the Yahoo tests, the network is organized in a two-tier tree topology with 5 machines connected to one Gigabit Ethernet leaf switch, 11 machines connected to another equal one, and both interconnected by a core switch through individual 10Gbps single links. Every machine has 8 processing cores provided by two Intel Xeon L5420 CPUs, each containing 4 cores and 1 thread per core. The available memory on each node is 8GiBs. On the Cervino tests, the network is organized in a flat topology with all nodes connected directly to the same switch via 40Gbps individual links. Every machine has 32 processing cores provided by two Intel Xeon E5-2683 v4 CPUs, each containing 16 cores and 2 threads per core. The available memory on each node is 128GiBs. For both experiments the employed operating system was Ubuntu 20.04.2; the chosen MPI implementation was Open MPI 4.0.3; and the benchmark employed was the OSU Micro benchmarks on version 5.7, executing the collective test `osu_allgather`. Although OSU is an opaque benchmark [19], we have chosen it due to its wide utilization on MPI communication research.

##### A. Metrics and configurations

All metrics were collected through the OSU reports, namely minimum, maximum and average time of the Allgather call over all runs. The Sparbit algorithm for `MPI_Allgather` was developed in C and bundled into an Open MPI collective communication component. The compared algorithms were Bruck, Recursive Doubling, Ring and Neighbor Exchange. Each algorithm was selected for execution through command line arguments passed to the `mpirun` executable, with native algorithms chosen through the dynamic rules of the `tuned` component, and Sparbit selected by raising its component priority. Block sizes were varied following OSU's default limits and variation, starting at 1B and raising up to 1MiB on successive multiplications by 2, totalling 24 tested data sizes.

The maximum number of processes utilized on the Yahoo tests was 256, while on Cervino it was 320, respecting a limit of 2 processes per physical core of each machine for both infrastructures. The processes were varied following two arithmetic progressions, both with step difference of 8, the first starting at 8 and ending at 256 for Yahoo (320 for Cervino) to represent cases of even and power of two numbers of processes. The second starting at 5 and ending at 253 (317) represented cases of odd numbers of processes, where there is asymmetry on the algorithms' execution and on the mapping of processes to the machines. On total there were 64 (80) different numbers of processes, resulting in 1344 (1680) test cases from combinations with data sizes. Each singular test of the previous parameters was executed employing both sequential and cyclic mappings, allowing the

empirical validation of the discussion on Section III-C. In order to reduce the probability and effect of any unexpected machine or network events on the time of a test, all algorithms were sequentially executed with both mappings for each number of processes. Finally, each test was executed 50 times for statistical significance, with 5 additional warm-up executions to reduce the effects of any possible start-up skews.

#### V. EXPERIMENTAL ANALYSIS

On this section we present the results obtained from the experiments and the analysis of their data. Initially, four hybrid heat maps with a broad view of the experiments are provided on Figure 5, two for Yahoo and two for Cervino, showing sequential and cyclic cases. Every cell of each heat map indicates a particular test configuration, comprising a unique combination of block size and number of processes. Each row indicates one block size from 1B to 1MiB and each column indicates one process amount, with both arithmetic progressions intertwined and starting from the odd one (as discussed on Section IV-A). The heat maps are said to be hybrid because they present both categorical and continuous values: if Sparbit does not have the smallest time on a test, the color of the cell indicates what traditional algorithm was the best (*e.g.*, the dark green color for 1B and 5 processes on Figure 5b indicates that Bruck was the best on this configuration). In turn, if Sparbit has the smallest time, a lighter is better grey scale color indicates its improvement percentage over the second best algorithm – *e.g.*, 23.83% improvement for 1MiB and 320 processes on Figure 5c. All the heat map values are related to the smallest average time over 50 executions.

**Network topology and data locality impact:** Most of the results that can be seen on Figure 5 corroborate with the previous discussions on Hockney-based algorithm costs, presented on Section II-A. Under sequential mapping on both infrastructures, Sparbit has the smallest execution time on a large portion of the tests, reaching 46.43% on Yahoo and 39.64% on Cervino, as seen on Figures 5a and 5c, respectively. Under cyclic mapping there is a considerable drop on Sparbit's percentage, as expected from its worse communication pattern on this case, reaching 19.12% on Yahoo and 30.83% on Cervino, as seen on Figures 5b and 5d, respectively. The different topologies of the infrastructures give valuable insights on its impact over Sparbit and the results itself. First, Yahoo has a larger topology with more levels crossing switches, which implies on different communication costs among distinct pairs of machines. Cervino on the other hand has a flat topology, which implies on equal theoretical communication costs among any two machines, leaving only the internal machine communication hierarchy and its difference to the external one. Therefore, on Yahoo with sequential mapping we can expect the highest benefit from employing Sparbit, which can be seen true by the largest fraction of the best times achieved on this

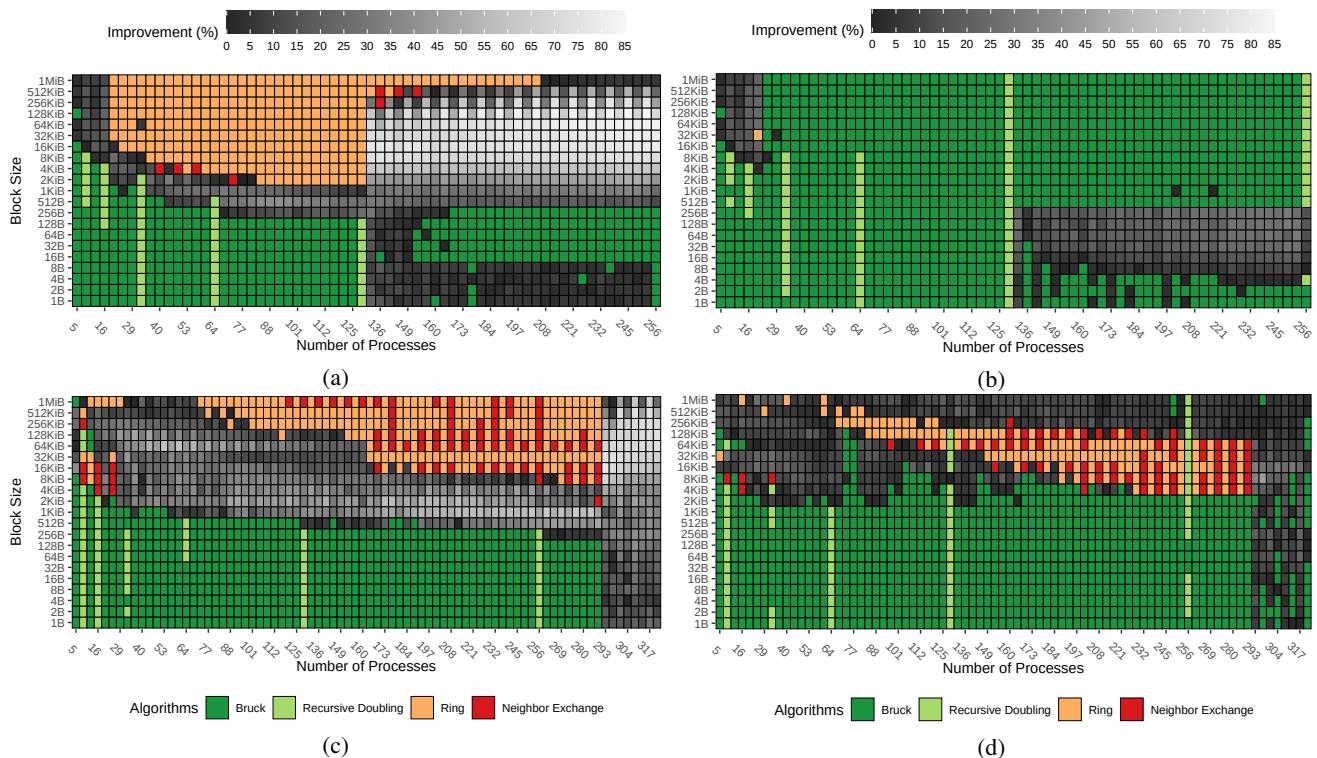


Figure 5: Heat maps of best average time algorithm for each case. (a) Yahoo cluster with sequential mapping. (b) Yahoo cluster with cyclic mapping. (c) Cervino cluster with sequential mapping. (d) Cervino cluster with cyclic mapping.

Cluster	Mapping (Total)	Min	Avg	Max	Min ∩ Avg	Min ∩ Max	Avg ∩ Max	Min ∩ Avg ∩ Max
Yahoo	Sequential (746)	115 (15.42%)	5 (0.67%)	6 (0.8%)	113 (15.15%)	1 (0.13%)	14 (1.88%)	492 (65.95%)
	Cyclic (346)	48 (13.87%)	0 (0%)	35 (10.12%)	5 (1.45%)	6 (1.73%)	33 (9.54%)	219 (63.29%)
Cervino	Sequential (799)	133 (16.65%)	6 (0.75%)	0 (0%)	196 (24.53%)	0 (0%)	3 (0.38%)	461 (57.7%)
	Cyclic (738)	202 (27.37%)	20 (2.71%)	16 (2.17%)	92 (12.47%)	2 (0.27%)	115 (15.58%)	291 (39.43%)

Table I: Relation of Sparbit's best minimum, average and maximum time sets.

experiment. The smaller topology of Cervino reduces the maximum potential benefit of Sparbit, which correspondingly reduces its fraction of the best times. Under cyclic mapping, Sparbit's drop in performance is expected to be proportional to the topology's size, since now there is an inverted communication pattern. This observation justifies the drop in relation to the sequential mapping seen on the Yahoo cluster to be larger than the one seen on Cervino, as the negative impact of the former's topology is higher than the latter. From all this we can state that the larger the topology's hierarchy, the larger the potential benefit of using Sparbit under sequential mapping, and the larger its benefit's reduction when switching to cyclic.

**Segmented analysis of algorithms' relation:** The form of the area taken by the Sparbit's best cases (Figure 5), on the particular intervals where it is the best, assume these uncommon shapes due to the cropped nature of the heatmap. If we individually compare the area of smaller groups of algorithms with Sparbit, like the ones with linear or logarithmic complexity, then more comprehensible forms

appear. Through this rationale, the forms of Sparbit on the Yahoo and Cervino sequential heat maps are the product of the same phenomenon. Considering Yahoo with sequential mapping (Figure 5a) and stripping Bruck and Recursive Doubling out of the plot, their area would be almost completely taken by Sparbit, absorbing 95.48% of the cases as the best, with Ring and Neighbor Exchange absorbing only the remaining percentage. If we instead stripped out Ring and Neighbor exchange while maintaining the rest, a similar pattern would emerge with Sparbit taking 100% of their area as the best. These observations show how there is a clear improvement pattern of Sparbit towards each of the algorithm groups. In relation to Ring and Neighbor Exchange, Sparbit is better on large numbers of processes and large data sizes (hereafter considered the ones larger than 1KiB), while it is better on almost all tests (97.44%) for smaller sizes (hereafter considered the ones smaller or equal than 1KiB). In relation to Bruck and Recursive Doubling, Sparbit is better for large data sizes also on almost all tests (97.16%). Therefore on the merged final version, which is

shown on Figure 5a, for 128 processes or less Sparbit excels on the verge sweet spot of both linear and logarithmic class algorithms, where the data size is too small to compensate the additional steps of the former and too large to compensate the non-local communication of the latter. Above 128 processes, which can be seen as a behavioural turning point, Sparbit takes over the majority of the large data sizes as the best - since it is the case on both the stripped versions - and a fraction of the small ones as well.

The same general pattern of Yahoo under sequential mapping can also be seen on Cervino. The difference is that the verge sweet spot for each number of processes is larger (*i.e.*, involves more test cases) and the behavioural turning point appears on 288 processes. We assume that these abrupt changes of behaviour are due to the machines entering an overbooked scenario where there are more process preemptions. On Yahoo, the first load where there are more processes than cores is with 133, and Cervino by having two threads can suppress this problem until higher ratios, but also fails before reaching 2 processes per core. This negatively affects all algorithms but is less severe on Sparbit, particularly on large data sizes, suggesting that it has less trouble when operating on highly utilized or restricted CPU environments.

Under cyclic mapping, Ring and Neighbor Exchange also suffer a drop in performance, since their neighbor ranks on the sequential mapping are now placed at other possibly far machines, which results in a very non-local communication pattern. Regarding Sparbit, on Cervino with cyclic mapping the area where it is the best results from a similar merging phenomenon as discussed on the sequential mapping. Here, the difference is that it is better on more cases of large data sizes previously taken by Ring and Neighbor Exchange, which are surpassed due to their more expressive drop on performance. On the other hand, it also loses more cases of both small and large data sizes to Bruck, due to the latter's improved communication locality on this mapping. On Yahoo with cyclic mapping, the form taken by Sparbit as the best virtually becomes a product of its relation with Bruck and Recursive Doubling, as it would on some of the previously discussed stripped versions of the plot. Sparbit is the best on the small top left area due to its use of parallel sends, employing more available bandwidth sooner, thus resulting in smaller communication time.

**Smallest minimum and maximum times:** The additional metrics for smallest minimum and maximum times are discussed in relation to the previous considerations. The forms taken by Sparbit as the best on the minimum and maximum heat maps are very similar to the ones taken on the already shown plots of average. Therefore, instead of presenting new heat maps we recur to Table I, which presents the relations amongst the minimum, average and maximum sets comprising cases where Sparbit has the best time.

For each infrastructure, Table I presents each mapping with the union of all cases in which Sparbit was the best.

The amount of cases with exclusively minimum, average or maximum smallest time are also presented on the third, forth and fifth columns, respectively. From columns 6 to 8 the binary set intersections are shown and column 9 indicates the cases on which Sparbit is simultaneously the best on the three sets. Initially, focusing only on the sequential mapping, Sparbit has either the minimum, average or maximum best time on 746 (55.51%) cases for Yahoo and 799 (47.56%) for Cervino. An important behaviour to notice is that for both infrastructures on this mapping there is little to no cases of exclusively average or maximum best times. This indicates that these metrics are correlated and that, as can be seen on the last two columns, on the majority of cases where it has the smallest maximum time, it also has the smallest average time – 98.64% (506/513) on Yahoo and 100% on Cervino. On both infrastructures the largest set of Sparbit's best times is the minimum, which also accounts for the largest exclusive number of cases and individual intersections with the average. This indicates that under sequential mapping it is easier for Sparbit to lower the minimum and average times than it is to lower the maximum ones. Such observation is further based on the decreasing amount of best cases from minimum to average and then maximum sets. The intersection of the three sets accounts for 36.61% (492/1344) of the total tests on Yahoo and 27.44% (461/1680) on Cervino, indicating that on these fractions it performs better than the other algorithms in a consistent way.

Under cyclic mapping, the total number of cases where Sparbit is the best in any of the metrics drops to 346 (25.74%) for Yahoo and 738 (43.93%) for Cervino, which follows the decay pattern seen previously in detail for the average. Here, the cases are more evenly spread along exclusive and binary intersections, especially on Cervino. The correlation between average and maximum metrics still exists but weaker, with 86% (252/293) and 95.75% (406/424) of the maximum best cases also being the best average ones for Yahoo and Cervino, respectively. On the latter, the minimum exclusive cases are very expressive, probably due to the fastest and more capable machines which amplify Sparbit's ability to reduce the minimum metric. There is no clear ordering of set sizes that applies to both infrastructures under this mapping and the intersection of all metrics drops to 16.29% (219/1344) of the total cases for Yahoo and 17.32% (291/1680) for Cervino.

**Sparbit's time reduction:** So far, the analysis has only focused on the number of cases in which Sparbit performed better than other algorithms, without discussing the actual improvement obtained on the tests. Table II summarizes the mean, median and highest percentage improvement values obtained on all three measured metrics for the different experimental scenarios. We also refer back to the heat maps of Figure 5 in order to discuss the patterns of improvement for the proposal. The first clear characteristic, as expected, is that the improvements under sequential mapping are much more significant and consistent than the ones under

Mapping	Cluster	Metric	Mean	Median	Highest
Sequential	Yahoo	Min	36.18	29.98	85.87
		Avg	34.70	26.16	84.16
		Max	35.48	27.24	80.85
	Cervino	Min	33.08	31.69	80.20
		Avg	30.23	29.00	77.78
		Max	27.67	26.14	74.27
Cyclic	Yahoo	Min	13.35	12.38	42.31
		Avg	14.89	15.77	31.07
		Max	16.07	16.64	36.47
	Cervino	Min	13.49	11.99	52.66
		Avg	09.60	08.71	44.12
		Max	09.71	08.43	37.36

Table II: Sparbit percentage improvement of metrics over the second best algorithm.

cyclic mapping. This is outlined on Figure 5 by a greater occurrence of darker tones for the plots of the latter on both infrastructures. The same phenomenon can also be seen numerically on Table II, as all the metrics present the smallest relative values on the cyclic side. Focusing on the plots for sequential mapping, another interesting feature is that the highest improvements appear generally on larger number of processes and data sizes. For Yahoo, above 128 processes and between 1KiB and 256KiB (287 tests), the mean (median) improvements are 60.64% (69.16%). On Cervino, above 168 processes and on the same data size interval, the mean (median) improvements are 43.86% (45.65%). The fact that on such cases the median is higher than the mean indicates a distribution that is asymmetric to the left, and thus has more larger improvements than lower ones. This difference is also higher on Yahoo, which supports the potential for growing benefits on larger topologies.

On a general picture the improvements are smaller but still very significant, as can be seen on Table II. Under sequential mapping Yahoo and Cervino values are quite similar, with means for the average time improvement respectively residing on about 35% and 30%, highest near 80% on both scenarios and average minimum as well as maximum times following similar trends. Under cyclic mapping, the improvements for average time are reduced on both infrastructures, achieving around 15% on Yahoo and 10% on Cervino. On this case, minimum and maximum metrics also follow similar behaviours, but the highest obtained improvements are more variable across all measurements ranging from 31% to 52%.

## VI. RELATED WORK

The first set of related work is composed by the already presented and compared classic Allgather algorithms (Section II), namely Bruck [10], Neighbor Exchange [11], Ring and Recursive Doubling. Another form of improving the time of the Allgather operation, specially regarding communication locality, is through rank reordering or remapping. On this approach the ranks assigned to a machine are swapped by others in order to make the application’s communication pattern better suited to the topology and

machine’s hierarchy. One very direct form of this technique is the one presented by [7], which changes the mapping from sequential to cyclic and vice-versa depending on the algorithm being executed. This approach shows great results for converting adverse mapping scenarios to good ones, with only local transformation functions and very little overhead. However, on the current form it requires gathering infrastructure information and supposes only equal distribution of processes to machines, thus on highly heterogeneous environments or with irregular mappings it would be unable to deliver great improvements.

The works of [16] and [6] have a more Allgather focused approach, using its known communication pattern to create mappings more suited for the algorithms. The first proposes fine-tuned heuristics for Ring, Recursive Doubling and Binomial broadcast (a possible final component of an Allgather or Broadcast execution), with the experimental results presenting improvements up to 78%. It needs however, either initial sends or final memory shifting to correct the blocks’ positions, which adds overhead and even negatively degrades the performance. The second work employs graph partitioning and linear optimization to fit the communication pattern to the topology, presenting high potential improvements but requiring infrastructure information and additional computation time, which grows for large numbers of processes.

The works of [20] and [21] focus on mappings for hierarchical tree networks and respectively propose the TreeMatch and Isomorphic Tree Mapping algorithms. These approaches assume that the topology structure is a balanced tree and require the profiling of applications in order to extract communication patterns. The information is then employed to map the ranks via the introduced algorithms. Both proposals show improvements over default forms of mapping, however require profiling information. In turn, the work of [22] proposes online monitoring and rank remapping that provide improvements and does not need prior executions, however it still requires active modification of application code. Sparbit could be potentially coupled with these techniques, however its main advantage in comparison is that it works out of the box, providing significant improvements on communication time for theoretically any hierarchical network, and without need for topology information, additional communication or computation.

## VII. CONCLUSION

On this paper we have highlighted the existing dichotomy of the employed Allgather algorithms for MPI, in which the best theoretical ones have physical limitations, mainly regarding communication locality, while the ones that present more local exchanges have a poor theoretical complexity. To address this current situation, the Stripe Parallel Binomial Trees (Sparbit) algorithm has been developed and presented, utilizing binomial trees to deliver the data faster, with optimal costs, greater locality and no usage restrictions.

Sparbit was empirically compared against the available algorithms on Open MPI 4.0.3, namely Bruck, Recursive Doubling, Ring and Neighbor Exchange. The experiments were executed on two HPC infrastructures with different topologies, network speeds and machine configurations to enrich the data sources for the analysis. The results obtained indicate that Sparbit achieves the smallest minimum, average and maximum times on a large fraction of test cases on the two infrastructures. The sequential mapping is the one where Sparbit ameliorates more cases and does so more intensively, however under cyclic a considerable amount of tests have also been improved.

The implementation employed here and geared towards the original MPI Allgather operation served as a proof of concept for the potential benefits of the algorithm. Nonetheless, it can be easily implemented to execute the Vector version of the call as well as the non-blocking forms, which we plan to develop and benchmark as future work. Upon this expansion, Sparbit can be used on several MPI collective operations.

**Acknowledgement:** This study was financed in part by the Coordenação de Aperfeiçoamento de Pessoal de Nível Superior - Brasil (CAPES) - Finance Code 001. It was also supported by FAPESC, UDESC, LabP2D, and University of Neuchâtel.

## REFERENCES

- [1] M. Forum, “MPI: A Message-Passing Interface Standard,” p. 868, 2015.
- [2] S. Chunduri, S. Parker, P. Balaji, K. Harms, and K. Kumaran, “Characterization of MPI Usage on a Production Supercomputer,” in *SC18: Int. Conf. for High Performance Computing, Networking, Storage and Analysis*, Nov. 2018, pp. 386–400.
- [3] Y. Gong, B. He, and J. Zhong, “Network Performance Aware MPI Collective Communication Operations in the Cloud,” *IEEE Transactions on Parallel and Distributed Systems*, vol. 26, no. 11, pp. 3079–3089, Nov. 2015.
- [4] I. Laguna, R. Marshall, K. Mohror, M. Ruefenacht, A. Skjellum, and N. Sultana, “A large-scale study of MPI usage in open-source HPC applications,” in *Int. Conf. for High Performance Computing, Networking, Storage and Analysis*. ACM, Nov. 2019, pp. 1–14.
- [5] G. Bosilca, C. Foyer, E. Jeannot, G. Mercier, and G. Papaure, “Online Dynamic Monitoring of MPI Communications,” in *Euro-Par 2017: Parallel Processing*, ser. Lecture Notes in Computer Science, F. F. Rivera, T. F. Pena, and J. C. Cabaleiro, Eds. Cham: Springer International Publishing, 2017, pp. 49–62.
- [6] M. G. Kurnosov, “Dynamic mapping of all-to-all collective operations into hierarchical computer clusters,” in *13th Int. Scientific-Technical Conf. on Actual Problems of Electronics Instrument Engineering*, vol. 02, Oct. 2016, pp. 475–478.
- [7] J. M. Alvarez-Llorente, J. C. Diaz-Martin, and J. A. Rico-Gallego, “Formal Modeling and Performance Evaluation of a Run-Time Rank Remapping Technique in Broadcast, All-gather and Allreduce MPI Collective Operations,” in *17th IEEE/ACM Int. Symp. on Cluster, Cloud and Grid Computing*, May 2017, pp. 963–972.
- [8] G. D. Benson, C.-W. Chu, Q. Huang, and S. G. Caglar, “A Comparison of MPICH Allgather Algorithms on Switched Networks,” in *Recent Advances in Parallel Virtual Machine and Message Passing Interface*, ser. Lecture Notes in Computer Science, J. Dongarra, D. Laforenza, and S. Orlando, Eds. Springer, 2003, pp. 335–343.
- [9] S. Kumar, S. S. Sharkawi, and K. A. N. Jan, “Optimization and Analysis of MPI Collective Communication on Fat-Tree Networks,” in *2016 IEEE Int. Parallel and Distributed Processing Symp. (IPDPS)*, May 2016, pp. 1031–1040, iSSN: 1530-2075.
- [10] J. Bruck, Ching-Tien Ho, S. Kipnis, E. Upfal, and D. Weathersby, “Efficient algorithms for all-to-all communications in multiport message-passing systems,” *IEEE Transactions on Parallel and Distributed Systems*, vol. 8, no. 11, pp. 1143–1156, Nov. 1997.
- [11] Jing Chen, Linbo Zhang, Yunquan Zhang, and Wei Yuan, “Performance evaluation of Allgather algorithms on terascale Linux cluster with fast Ethernet,” in *Eighth Int. Conf. on High-Performance Computing in Asia-Pacific Region*, Nov. 2005, pp. 6 pp.–442.
- [12] R. Thakur, R. Rabenseifner, and W. Gropp, “Optimization of Collective Communication Operations in MPICH,” *The International Journal of High Performance Computing Applications*, vol. 19, no. 1, pp. 49–66, Feb. 2005.
- [13] A. Faraj, X. Yuan, and D. Lowenthal, “STAR-MPI: self tuned adaptive routines for MPI collective operations,” in *20th Annual Int. Conf. on Supercomputing*, ser. ICS ’06. New York, NY, USA: Association for Computing Machinery, Jun. 2006, pp. 199–208.
- [14] S. Hunold and A. Carpen-Amarie, “Algorithm Selection of MPI Collectives Using Machine Learning Techniques,” in *2018 IEEE/ACM Performance Modeling, Benchmarking and Simulation of High Performance Computer Systems (PMBS)*, Nov. 2018, pp. 45–50.
- [15] R. W. Hockney, “The communication challenge for MPP: Intel Paragon and Meiko CS-2,” *Parallel Computing*, vol. 20, no. 3, pp. 389–398, Mar. 1994.
- [16] S. H. Mirsadeghi and A. Afsahi, “Topology-Aware Rank Reordering for MPI Collectives,” in *IEEE Int. Parallel and Distributed Processing Symp. Workshops*, May 2016, pp. 1759–1768.
- [17] J. Wang, D. Wang, M. Qiu, Y. Chen, and B. Guo, “A locality-aware shuffle optimization on fat-tree data centers,” *Future Generation Computer Systems*, vol. 89, pp. 31–43, 2018.
- [18] T. Angskun, G. Bosilca, and J. Dongarra, “Binomial graph: A scalable and fault-tolerant logical network topology,” in *Parallel and Distributed Processing and Applications*, I. Stojmenovic, R. K. Thulasiram, L. T. Yang, W. Jia, M. Guo, and R. F. de Mello, Eds. Berlin, Heidelberg: Springer Berlin Heidelberg, 2007, pp. 471–482.
- [19] L. Stanisic, L. M. Schnorr, A. Degomme, F. C. Heinrich, A. Legrand, and B. Videau, “Characterizing the performance of modern architectures through opaque benchmarks: Pitfalls learned the hard way,” in *2017 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW)*, 2017, pp. 1588–1597.
- [20] E. Jeannot and G. Mercier, “Near-Optimal Placement of MPI Processes on Hierarchical NUMA Architectures,” in *Euro-Par 2010 - Parallel Processing*, ser. Lecture Notes in Computer Science, P. D’Ambra, M. Guarracino, and D. Talia, Eds. Berlin: Springer, 2010, pp. 199–210.
- [21] D. Li, Y. Wang, and W. Zhu, “Topology-Aware Process Mapping on Clusters Featuring NUMA and Hierarchical Network,” in *2013 IEEE 12th Int. Symp. on Parallel and*

- Distributed Computing*, Jun. 2013, pp. 74–81, iSSN: 2379-5352.
- [22] E. Jeannot and R. Sartori, “Improving MPI Application Communication Time with an Introspection Monitoring Library,” in *IEEE Int. Parallel and Distributed Processing Symp. Workshops*, May 2020, pp. 691–700.